

PaTAS latency overhead — inference (left) and training in-process (right, 5 epochs)

