


PaTAS latency overhead — inference (left) and training (right, 5 epochs)

Inference latency vs batch size



Training latency across noise conditions

